

# Transport Layer

- Transport Layer provides for logical communication b/w application processes running on different hosts
- Logical communication: As if hosts<sup>running processes</sup> were directly connected // direct, end-to-end connection
  - Provides communication services directly to the application processes
- Features:
  - Implemented in end systems
  - On sending side, it converts application-layer messages to into transport layer packets called segments
    - Done by breaking messages into smaller chunks and adding a transport-layer header to each chunk to create a segment
  - segment passed to network layer where it is encapsulated to form a datagram (network-layer packet) and sent to destination

## Note:

- Routers only act on network layer fields i.e: they do not examine transport-layer segment fields encapsulated with the datagram
- On receiving end, router extract the segment and passes it to transport layer
- Transport processes segment and passes the message to application layer
- Transport layer provides logical communication b/w processes
- Network Layer provides logical communication b/w hosts
- Certain services can be provided by transport layer if the underlying layer doesn't provide it

Q- 2 houses , very distant from each other . In respective houses Bill, Anne responsible for delivering and collecting mails from house members. Postal service responsible for delivering mails b/w houses

- Message= letters
- processes= house members
- Transport layer protocol = Anne, Bill
- Network layer protocol = Postal service
- Protocols at transport layer:
  - TCP (Transmission Control Protocol)
  - UDP ( User Datagram Protocol)
- Common features of UDP and TCP
- Both protocols extend host-to-host delivery to process-to-process delivery and is called transport layer multiplexing and de-multiplexing
- Provide integrity checking by including error-detection fields in their headers
- TCP
  - Provides reliable data transfer - flow control, seq numbers, congestion control, acks, timers
  - Ensures correct and in-order delivery to the receiving process.
  - congestion control prevents TCP connections from swamping the links and routers b/w communicating hosts with an excessive amount of traffic - done by giving the connection an equal share of the link bandwidth, by regulating rate of sending side at which they can send traffic into network
- UDP
  - Unreliable
  - UDP traffic is unregulated // can send data at any rate
- TCP converts IP's unreliability by providing features that IP doesn't i.e: guaranteed correct order, no corruption, acks ....

# Multiplexing

- A process can have one or more sockets

- De-multiplexing:  
At the receiving end, transport layer examines the segment headers // fields to identify the receiving socket, and directs the segment to that socket

collect

## Multiplexing:

- Gathering of data chunks at the source host from different sockets, encapsulating each of them with header information to create segments, and passing the segments to network layer

- Sockets have unique identifiers → Port Numbers
- Each segment has special fields

- source -Port and destination Port numbers = special Fields
- Each port num is 16-bit.
- Ports 0–1023 are well-known ports // reserved for protocols like HTTP ....

80

- Data chunks are stream of bytes that are received from sockets. After attaching header info, the segments are divided into packets and sent to network layer

## Connection - Connection Less Multiplexing and demultiplexing

- TCP socket as a 4-tuple model  
→ (sourceIP, source port, destinationIP, destination port)

- TCP connection-oriented
- UDP connection-less

- connection establishment request sent to server, containing connection establishment bit set in TCP header, source port and destination port number → socket (AF\_INET, SOCK\_STREAM)  
· connect(server, dest port)

- when the host OS running the server process receives the connection request, it locates the server process waiting to accept the connection on the port number and creates a new socket accept()

- The new socket created is identified by the 4-tuples in the connection request source IP, source port, destination IP, destination port

- UDP uses dest port only for demux
- TCP uses all 4 tuples for demux

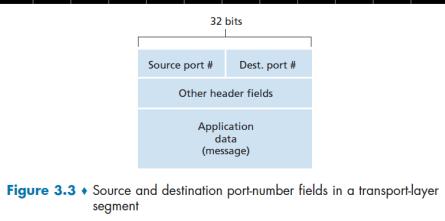
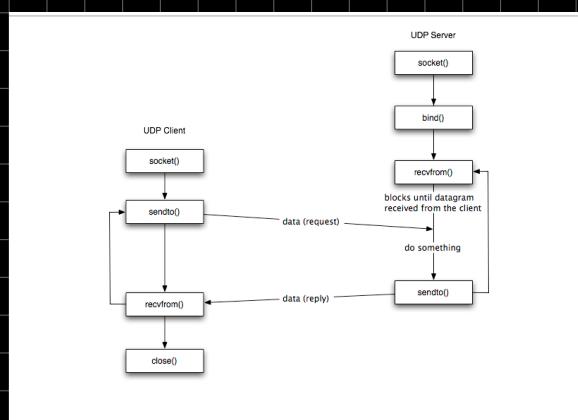
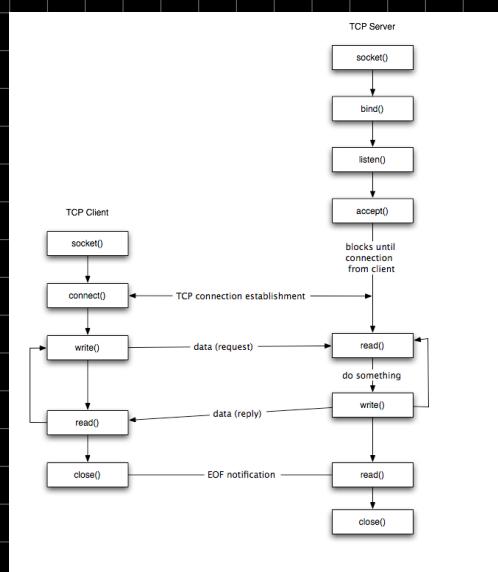


Figure 3.3 • Source and destination port-number fields in a transport-layer segment

- TCP: Ack for every message sent



## Socket workflows



UOP

- Connection less because no handshaking b/w client and receiver e.g: DNS uses UDP
  - Host waits for reply after sending, if no reply then resend with some changes e.g: send to another name server (not compulsory)
  - Why UDP is better ?

- Finer application-level control over what data is sent and when
    - Suited for real-time applications requiring minimum sending rate, do not want overly delay segment transmission, and can tolerate some data loss.
    - These additional services can be implemented as part of the application if required
  - No connection establishment
    - Eliminates delay for connection establishment (DNS)

- NO connection state
    - Allows to run more active clients
    - does not store receive, send buffers, congestion control and sequence, acknowledgement params

- Small packet overhead
    - UDP segment has 8 bytes of overhead
    - TCP segment has 20 bytes of overhead
  - Many applications switch to UDP like network management, DNS, new versions of HTTP, some streaming as well

# segment structure vOP

- length field specifies number of bytes in the segment (header + data)
  - UDP header has only 4 fields each of 2 bytes // 16 bit
  - checksum for error detection each

## UDP header

32 bits	
Source port #	Dest. port #
Length	Checksum
Application data (message)	

**Figure 3.7** ♦ UDP segment structure

- **Checksum:**  
Sender Side ↓
  - Source port and destination port is summed up
  - If any overflow, then it is added to the LSB side (wrap around)
  - 1's complement applied to the sum and added to checksum field

\* different in slides

Q- 0110011001100000  
0101010101010101  
1000111100001100

overflow

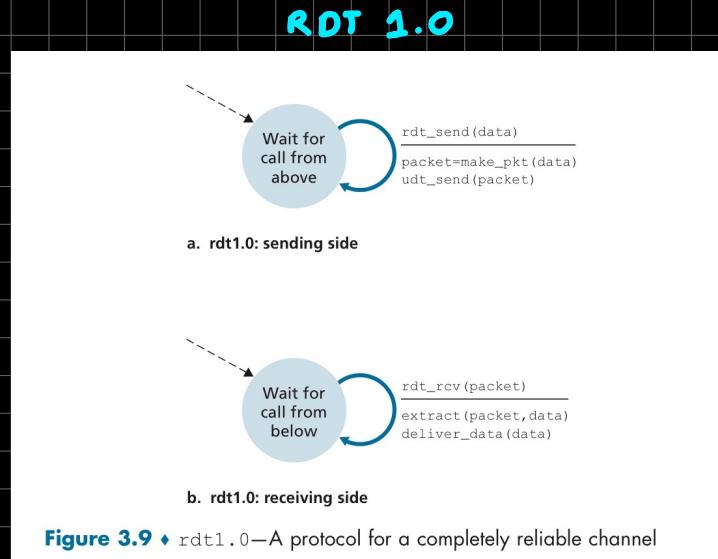
0100101011000010	$\begin{array}{r} 100011110000110 \\ \times 10000000000000001 \\ \hline 0100101011000010 \end{array}$
------------------	---

- Final String = 1011 0101 0011 1101
  - At receiver end, all four 16 bit words are added, and if no errors are introduced the checksum will result in 1111 1111 1111 1111.

A horizontal number line consisting of a black line with five evenly spaced tick marks. Two red 'x' marks are placed on the line, one at each end, indicating the range of the inequality.

# Reliable Data Transfer

- Responsibility of delivering data in-order, without any corruption and data losses
- Network layer is an unreliable channel, so implemented in transport layer
- RDT is a service abstraction that is made real by TCP



- rdt = reliable data transfer
- udt = unreliable data transfer
- Everything discussed is uni-directional

- Assuming the underlying channel (network layer) is completely reliable
- **Sending Side:**
  - `rdt.send(data)` event, creates a packet containing data using `make_pkt(data)`, and sends it to the channel using `udt.send()`
  - `rdt.send()` invoked by upper-layer application
- **Receiving Side:**
  - `rdt.recv(packet)` event, removes data from packet using `extract(packet, data)` and passes data to upper layer via `deliver_data(data)`
  - `rdt.recv()` invoked from underlying layer
- No need for receiver to ask sender to slow down or send any feedback

## ROT 2.0

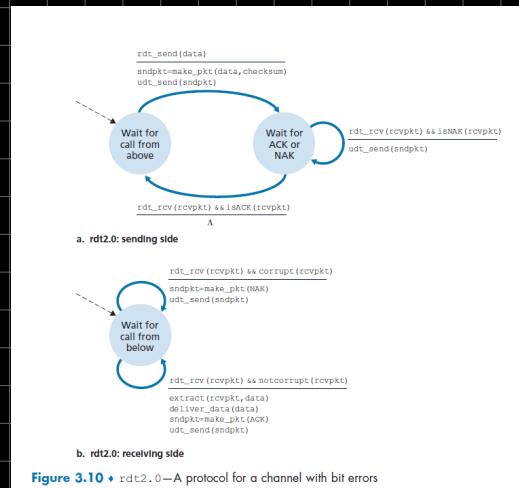


Figure 3.10 • rdt2.0—A protocol for a channel with bit errors

- Assuming packets sent are received (can have errors), and not lost
- uses positive acknowledgments (Ack) and negative acknowledgments (NAK)
- These control messages allow sender to know, that data has been received correctly/incorrectly
- Retransmission protocols known as automatic repeat request protocols (ARQ)
- ARQ features to handle bit errors:
  - Error Detection techniques e.g. checksum
  - Receiver feedback using Ack's and NAK's
  - Retransmission of packet that is received in error
- Sending side:
  - sndpkt contains data to be sent along with checksum
  - Ack or NAK is waited for
  - If a NAK is received, the protocol retransmits the last packet and waits for an Ack
  - If Ack is received, the protocol moves to the state of waiting for the packet from the upper layer
  - Can't get more data from upper layer until Ack is received, hence rdt 2.0 is a stop-and-wait protocol
- Receiving side:
  - Simple workflow see diagram
- Flaw:
  - Ack/NAK packet can be corrupted
  - Can not retransmit b/c it would introduce duplicate packets
  - Adding additional checksum bits to make a mathematically strong checksum value, only applicable for channels that can corrupt packets but not lose them
- Solution is to add a sequence number of the packets. Receiver uses sequence number field to identify retransmission. Sequence number moves forward in a modulo 2 arithmetic fashion.
- Due to assumption, Ack's / NAK's do not need to identify packets they are acknowledging, it is obviously of the most recent packet sent

## ROT 2.1

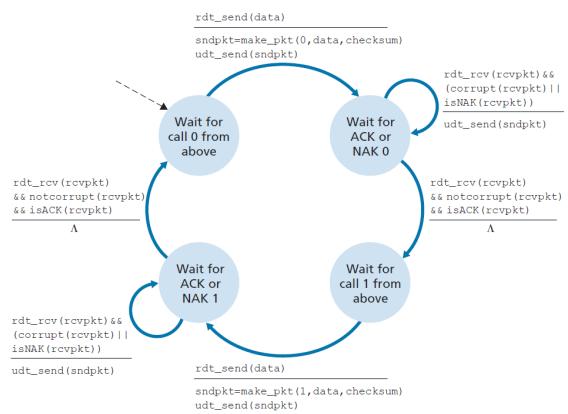


Figure 3.11 • rdt2.1 sender

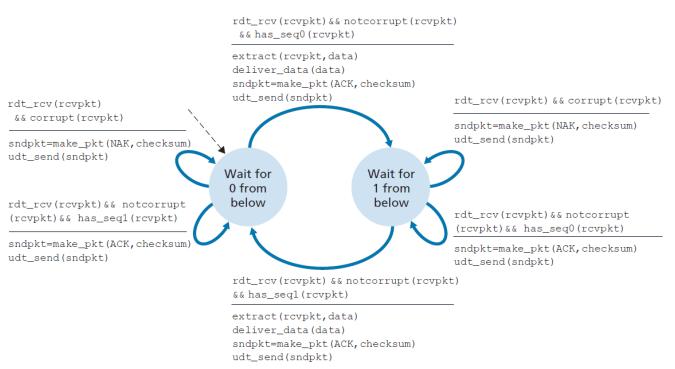
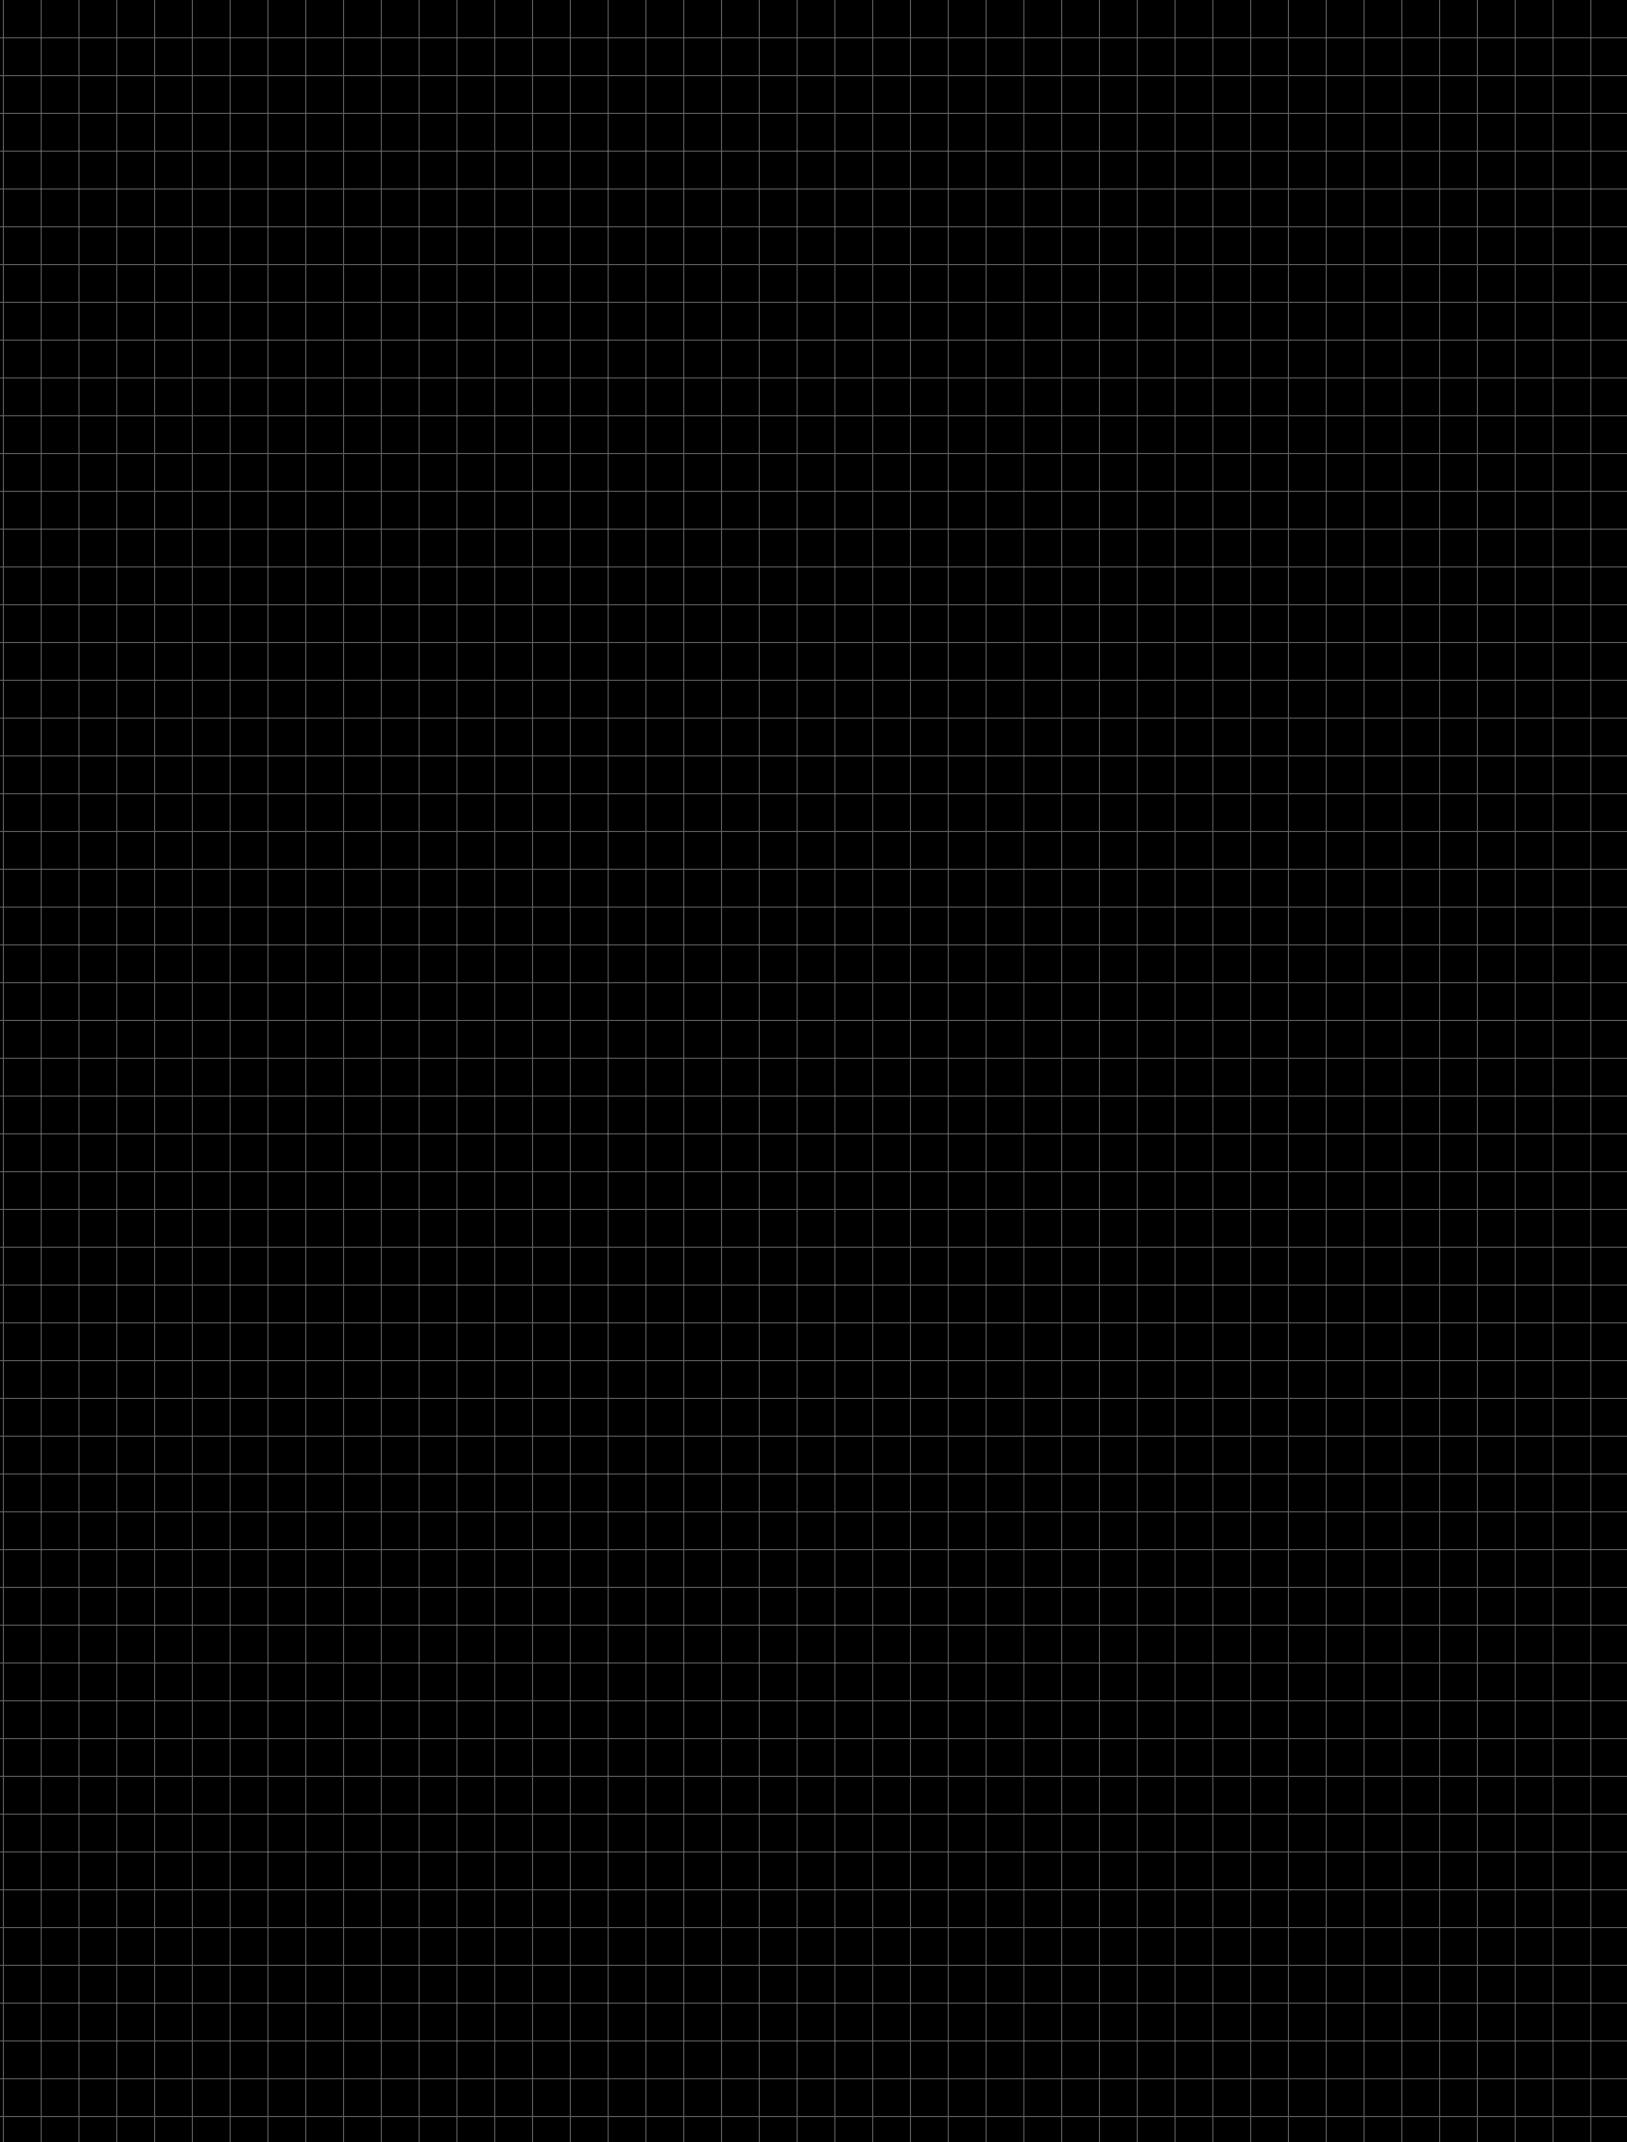


Figure 3.12 • rdt2.1 receiver

- when an out-of-order packet received, ACK is sent.
- when corrupted packet is received, NAK is sent
- A sender that receives two acks for the same packet (duplicate ACKs) means receiver did not correctly receive the packet following the packet that is ACKed twice
- Receiver now includes sequence number of the packet being acknowledged by adding argument in make\_pkt by ACK0, ACK1.
- Sender examines the sequence number of packet being acknowledged through isACK

## ROT 2.2



- Many link-layer protocols provide error checking capability, so why does UDP provide checksum?

→ No guarantee that all links b/w source and destination provide error checking, some might be using protocols that do not provide error checking

→ Bit errors can be introduced in router's memory

  - Link by link reliability, and in-memory error detection not guaranteed
  - UDP does not do anything to recover from an error. Some implementations of UDP either discard or pass the damaged segment to application layer  
  - Length is of 16 bits size, so theoretical max size of UDP packet =  $2^{16} - 1 = 65535$
  - Most physical layers have MTU (Max Transmission unit) to dictate how much data can be transmitted in 1 packet
  - If UDP exceeds this amount, the segment will be fragmented according to MTU.
  - Fragmentation increases likelihood of packet loss, as a single fragment loss results in the loss of entire original datagram

→ Hence, UDP packets are limited to very small sizes to avoid fragmentation

  - Path MTU: smallest MTU along the entire path b/w 2 IP hosts
  - Typical UDP size =  $1500 - 20 = 1480$ ; 20 bytes of UDP header
  - For IPv4: 65535 bytes for entire IPv4 packet - 16 bit
  - Data that can be fit into a UDP datagram sent using IPv4 =  $65535 - 20 - 8$   

$$\begin{array}{r} \text{IP header} \quad \text{UDP header} \\ \hline \end{array}$$
  
 $= 65507$
  - For IPv6: 65535 bytes - IPv6 payload limit - 16 bit
  - Data =  $65535 - 8 = 65527$

Checksum:

  - optional for IPv4, mandatory in IPv6
  - Use 0x0000 when need to include checksum field in the checksum algo.
  - In IPv4 → 0x0000, UDP checksum not in use
  - If checksum value comes out to 0x0000, change it to 0xFFFF
  - If checksum is in use, and on receiver's side if checksum not verified, packet is usually thrown away
  - Error detection codes: Reed-Solomon codes, Hamming codes
  - Demux key at IP layer

→ RDT