

# Class Notes


## Preamble

- (1) One-third of the semester is over!
- (2) The syllabus of the second sessional exam will be **chapter 3** and **chapter 4** from our textbook.
- (3) All exams are comprehensive, meaning everything we studied in the class starting lecture one will be included in the syllabus.
- (4) This middle part of the course is relatively dense and fast-paced. It is **mandatory** to read the relevant sections from the textbook along with the lectures.

## A quick revision

Let's start with a quick revision to build our context.

## Connectionless Transport UDP



### (1) Structure of UDP segment

- a. By convention RFC's show 32 bits in a row. We will follow that convention.
- b. Be careful which bit is the first bit and hence will be put first on the wire.
- c. Source and destination ports are of 16 bit each (possible values from 0 to 65535 i.e.  $2^{16} - 1$ )

- i. Port no 0 to 1023 are “well-known” ports and your application will not be able to use them on OS like Linux until your application has appropriate privileges

(2) Length is of 16 bits

- a. Theoretical max size of a UDP packet (UDP header + payload) = 65535
- b. Units of Length = Bytes
- c. But IPv4 which will carry a UDP packet as payload can not be larger than 65535. Smallest IPv4 header is of 20 Bytes. So  $65535 - 20 = 65515$  Bytes
- d. But most physical layer technologies have a Maximum Transmission Unit (MTU) that dictates how much data they can deliver in one packet.
- e. If a UDP packet exceeds the MTU of a network segment, it will be fragmented into smaller pieces by the IP layer. Fragmentation increases the likelihood of packet loss, as the loss of a single fragment results in the loss of the entire original datagram. For this reason, many applications that use UDP, such as DNS, often limit their packet sizes to a much smaller and safer value, like 512 bytes, to avoid fragmentation.
- f. Path MTU: Find the minimum MTU end-to-end and use that size to avoid data packet fragmentation at the IP layer.
- g. Typical UDP size =  $1500 - 20 = 1480$  Bytes

Common MTU values — compact table		
Technology / Link layer	Typical MTU (bytes)	Notes
Ethernet (IEEE 802.3)	1500	De-facto Internet LAN MTU (IP payload).
Ethernet — Jumbo frames	~9000 (often 9000–9216)	Vendor-dependent; must be supported end-to-end to be useful.
Wi-Fi (IEEE 802.11)	1500 (typical IP MTU)	802.11 MAC can carry larger MSDUs (up to 2304) but stacks usually use 1500 for compatibility.
PPPoE (DSL)	1492	Ethernet MTU (1500) minus 8 bytes PPPoE overhead — common WAN setting.
LTE / Cellular	~1500 (varies: often 1350–1500)	Carriers and tunneling can reduce MTU; path-dependent.
Bluetooth Classic (L2CAP)	~672 (stack-dependent)	Historic default; can be negotiated larger depending on implementation.
Bluetooth LE (ATT)	23 (default), commonly negotiated up to 247	ATT default MTU 23; modern stacks and LE Data Length Extension allow larger L2CAP payloads.
FDDI (historical)	4352	Older LAN tech — shown for reference.
ATM (AAL5)	~9180 (AAL5 maximum payload)	ATM cell network commonly used with AAL5 for IP; practical limits depend on encapsulation.
Loopback (software, e.g. lo on Linux)	65,536 (common default)	Host-local interface; very large since no physical link.

- h. IPv4 header size is variable due to some reasons we will see later. Minimum IPv4 header size is 20 Bytes. IPv6 header is fixed at 40 Bytes.

- i. For IPv4: 65,535 bytes (IP total length)–20 bytes (IP header)–8 bytes (UDP header)= 65,507 bytes of data
- j. For IPv6: 65,535 bytes (IPv6 payload limit)–8 bytes (UDP header)=65,527 bytes of data

#### Key Differences Summarized

Feature	IPv4	IPv6
<b>Field Name</b>	Total Length	Payload Length
<b>What's Included</b>	<b>Header + Payload</b>	<b>Payload Only</b>
<b>How to get total size</b>	The value is the total size	40 bytes + field value

i.

(3) Checksum

(4) UDP checksum is over (pseudo header) AND (UDP header) AND (UDP Payload)

- a. Optional in Ipv4, mandatory in Ipv6
- b. Use 0X0000 when need to include checksum field in the checksum algo
- c. In Ipv4, 0X0000 => UDP checksum not in use
- d. To disambiguate, if checksum values comes out to 0X0000, change it to 0X1111
- e. If checksum is in use, and on receiver side if checksum not verified, packet is usually thrown away.

(5) Pseudo headers. Why we need them?

#### IPv4 pseudo-header (concatenate these bytes, in order):

1. Source address (32 bits)
2. Destination address (32 bits)
3. Zero (8 bits) — one zero octet, value `0x00`
4. Protocol (8 bits) — e.g. `0x11` for UDP
5. UDP/TCP length (16 bits) — network byte order

So the bytes layout is:

`[src(4)] [dst(4)] [0x00] [protocol] [length(2)]`

#### IPv6 pseudo-header (concatenate these bytes, in order):

1. Source address (128 bits)
2. Destination address (128 bits)
3. Upper-layer packet length (32 bits) — length of header+payload
4. Zero (24 bits) — three zero octets `0x00 0x00 0x00`
5. Next Header (8 bits) — e.g. `0x11` for UDP

So the bytes layout is:

`[src(16)] [dst(16)] [length(4)] [0x00 0x00 0x00] [next-header]`

- (6) If link-layer provides data integrity why do it again at UDP level?
  - a. End-to-end principle
- (7) Some discussion on error detection and error correction codes
  - a. Hamming code
  - b. Reed-Solomon codes
  - c. ... many more
- (8) UDP checksum example from slides (Slide no 34)
- (9) Another note: Note the DeMux key at the IP layer!